

DenkowskiStanislawlab6

December 2, 2022

```
[ ]: !pip install pycocotools --user
```

```
Collecting pycocotools
  Downloading pycocotools-2.0.6.tar.gz (24 kB)
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Requirement already satisfied: matplotlib>=2.1.0 in
/opt/conda/lib/python3.7/site-packages (from pycocotools) (3.5.3)
Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages
(from pycocotools) (1.21.6)
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.7/site-
packages (from matplotlib>=2.1.0->pycocotools) (0.11.0)
Requirement already satisfied: packaging>=20.0 in /opt/conda/lib/python3.7/site-
packages (from matplotlib>=2.1.0->pycocotools) (21.3)
Requirement already satisfied: pyparsing>=2.2.1 in
/opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.0->pycocotools)
(3.0.9)
Requirement already satisfied: fonttools>=4.22.0 in
/opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.0->pycocotools)
(4.33.3)
Requirement already satisfied: pillow>=6.2.0 in /opt/conda/lib/python3.7/site-
packages (from matplotlib>=2.1.0->pycocotools) (9.1.1)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.0->pycocotools)
(2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in
/opt/conda/lib/python3.7/site-packages (from matplotlib>=2.1.0->pycocotools)
(1.4.3)
Requirement already satisfied: typing-extensions in
/opt/conda/lib/python3.7/site-packages (from
kiwisolver>=1.0.1->matplotlib>=2.1.0->pycocotools) (4.1.1)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-
packages (from python-dateutil>=2.7->matplotlib>=2.1.0->pycocotools) (1.15.0)
Building wheels for collected packages: pycocotools
  Building wheel for pycocotools (pyproject.toml) ... done
  Created wheel for pycocotools:
filename=pycocotools-2.0.6-cp37-cp37m-linux_x86_64.whl size=373761
```

```
sha256=5d9fa500d156cb59ea01a5c29b8c7b58963c3aa8d7d96dd0c0bec0e3e7d5f9d9
  Stored in directory: /root/.cache/pip/wheels/06/f6/f9/9cc49c6de8e3cf27dfddd91b
f46595a057141d4583a2adaf03
Successfully built pycocotools
Installing collected packages: pycocotools
Successfully installed pycocotools-2.0.6
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv
```

```
[ ]: import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import tensorflow as tf
import tensorflow_addons as tfa
import numpy as np
import os
import time
```

```
[ ]: (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.
    ↪cifar10.load_data()
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog',
    ↪'horse', 'ship', 'truck']
train_labels, test_labels = tf.squeeze(tf.one_hot(train_labels, len(classes))),
    ↪tf.squeeze(tf.one_hot(test_labels, len(classes)))
train_images, test_images = tf.convert_to_tensor(train_images), tf.
    ↪convert_to_tensor(test_images)
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 4s 0us/step
170508288/170498071 [=====] - 4s 0us/step
```

```
[ ]: augment = tf.keras.Sequential([
    tf.keras.layers.RandomFlip(mode='horizontal'),
    tf.keras.layers.RandomZoom((-0.1,0.1)),
    tf.keras.layers.Normalization()
])
augment.layers[-1].adapt(train_images)
```

```
2022-12-02 14:22:48.567276: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR
Optimization Passes are enabled (registered 2)
```

Wczytałem Cifar10, następnie przygotowałem model, który wykonuje augmentację w postaci RandomFlip, RandomZoom oraz normalizację, ponieważ niestety tensorflow nie ma RandomResizedCrop a rozwiązania jakie znalazłem nie działały.

W tensorflow, z tego co rozumiem to albo wrzuca się przez to dane, albo podcina na początek

modelu i tak spróbuję(choć ostatnio miałem problem).

Dodatkowo zaaplikowałem one hot encoding na labele i zrzutowałem obrazki na tensor.

```
[ ]: class Patches(tf.keras.layers.Layer):
    def __init__(self, size):
        super().__init__()
        self.size = size



    def call(self, images):
        batch_size = tf.shape(images)[0]
        patches = tf.image.extract_patches(
            images=images,
            sizes=[1,self.size,self.size,1],
            strides=[1,self.size,self.size,1],
            rates=[1,1,1,1],
            padding='VALID'
        )
        patch_dims = patches.shape[-1]
        patches = tf.reshape(patches, [batch_size, -1, patch_dims])
        return patches
```

Klasa Patches zawiera implementację rozkładania obrazka na patche, wykorzystuję do tego `tf.image.extract_patches`.

Warto zauważyć, że klasa dziedziczy z `tf.keras.layers.Layer`, dzięki czemu może być częścią modelu.

```
[ ]: patches = Patches(4)
plt.imshow(train_images[0])
plt.show()
patched = patches(tf.expand_dims(train_images[0],0))

plt.figure(figsize=(4,4))
for ix, im in enumerate(patched[0]):
    plt.subplot(8, 8, ix + 1)
    im = tf.reshape(im, (4,4,3))
    plt.imshow(im)
    plt.axis('off')
```



```
[ ]: class Encoder(tf.keras.layers.Layer):
    def __init__(self, num_patches=8*8, projection_dim=256):
        super().__init__()
        self.num_patches = num_patches
        self.projection_dim = projection_dim
```

```

        init = tf.random_normal_initializer()
        class_token = init(shape=(1,self.projection_dim), dtype=tf.float32)
        self.class_token = tf.Variable(initial_value=class_token,
        ↪ trainable=True)
        self.projection = tf.keras.layers.Dense(units=self.projection_dim)
        self.positional_embedding = tf.keras.layers.Embedding(input_dim=self.
        ↪ num_patches+1, output_dim=self.projection_dim)

    def call(self, batch):
        batch_size = tf.shape(batch)[0]
        class_token = tf.tile(input=self.class_token, multiples=[batch_size,1])
        class_token = tf.reshape(class_token, shape=(batch_size,1,self.
        ↪ projection_dim))
        embs = self.projection(batch)
        embs = tf.concat([class_token, embs], 1)
        positions = tf.range(0, self.num_patches+1, 1)
        pos_embs = self.positional_embedding(positions)
        encoded = embs+pos_embs
        return encoded

```

Do stworzenia encodera użyłem: - `tf.random_normal_initializer` - do zainicjowania class tokena, inicjuję go trochę jako batch, dlatego wymiar (1,rozmiar tokena) - `tf.Variable` - aby class token się uczył i był używany przez framework - `tf.keras.layers.Dense` - do liniowego embeddingu patchy - `tf.keras.layers.Embedding` - do embeddingu pozycyjnego

Ponadto do wykorzystania go używam: - `tf.tile` - aby dodać `class_token` do każdego obrazka w batchu - `tf.reshape` - aby batch `class_tokenów` miał właściwe wymiary - warstwy dense, wcześniej zainicjowanej - do liniowego embeddingu patchy - `tf.concat` - aby dokleić do batcha `class_token`y - `tf.range` - aby wyznaczyć embeddingi pozycyjne - warstwy embedding, wcześniej zainicjowanej - do wyznaczenia embeddingu pozycyjnego - ona się będzie uczyć, ale potrzebuje dostać `tf.range` - sumuję embeddingi i otrzymuję encoding

```

[ ]: class Transformer(tf.keras.layers.Layer):
    def __init__(self, num_head=8, projection_dim=256, hidden_units=512):
        super().__init__()
        self.num_head = num_head
        self.projection_dim = projection_dim
        self.hidden_units = hidden_units
        self.norm1 = tf.keras.layers.LayerNormalization()
        self.multihead = tf.keras.layers.MultiHeadAttention(num_heads=self.
        ↪ num_head, key_dim=self.projection_dim)
        self.norm2 = tf.keras.layers.LayerNormalization()
        self.linear1 = tf.keras.layers.Dense(units=self.hidden_units,
        ↪ activation=tf.nn.gelu)
        self.dropout1 = tf.keras.layers.Dropout(rate=0.2)
        self.linear2 = tf.keras.layers.Dense(units=self.projection_dim)
        self.dropout2 = tf.keras.layers.Dropout(rate=0.2)

```

```

def call(self, batch):
    batch1 = self.norm1(batch)
    attentioned = self.multihead(batch1, batch1)
    self.attention_output = attentioned
    batch2 = tf.keras.layers.Add()([attentioned, batch])
    batch3 = self.norm2(batch2)
    batch3 = self.linear1(batch3)
    batch3 = self.dropout1(batch3)
    batch3 = self.linear2(batch3)
    batch3 = self.dropout2(batch3)
    batch4 = tf.keras.layers.Add()([batch3, batch2])
    return batch4

```

Aby stworzyć blok transformera użyłem: - `tf.keras.layers.LayerNormalization` - jako warstwy normalizującej - `tf.keras.layers.MultiHeadAttention` - jako warstwy multi head attention - `tf.keras.layers.Dense` - Jako liniowej warstwy - `tf.keras.layers.Dropout` - Jako warstwy dropout - `tf.nn.gelu` - jako funkcji aktywacji

Do wykorzystania użyłem jeszcze: - `tf.keras.layers.Add` - jako połączenie rezydualne

```

[ ]: def create_vit(patch_size=4, transformers_num=6):
    inputs = tf.keras.layers.Input(shape=(32,32,3))
    # augmented = augment(inputs)
    patches = Patches(size=patch_size)(inputs) #(augmented)
    embedded = Encoder()(patches)
    transformed = tf.keras.layers.Dropout(0.2)(embedded)
    for _ in range(transformers_num):
        transformed = Transformer()(transformed)
    normed = tf.keras.layers.LayerNormalization()(transformed[:,0])
    mlped = tf.keras.layers.Dense(units=len(classes),
    ↪activation='softmax')(normed)
    model = tf.keras.Model(inputs=inputs, outputs=mlped)
    return model

vit = create_vit()

```

Aby zdefiniować model, ustalam input, przepuszczam input przez augmentację, następnie rozkładam obrazki na patche, embedduje, nakładam bloki transformerów 6 razy, a na koniec jedna warstwa liniowa.

```

[ ]: vit.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
patches_1 (Patches)	(None, None, 48)	0

```

-----
encoder (Encoder)          (None, 65, 256)          29440
-----
dropout (Dropout)          (None, 65, 256)           0
-----
transformer (Transformer)   (None, 65, 256)         2367488
-----
transformer_1 (Transformer) (None, 65, 256)         2367488
-----
transformer_2 (Transformer) (None, 65, 256)         2367488
-----
transformer_3 (Transformer) (None, 65, 256)         2367488
-----
transformer_4 (Transformer) (None, 65, 256)         2367488
-----
transformer_5 (Transformer) (None, 65, 256)         2367488
-----
tf.__operators__.getitem (S1 (None, 256)          0
-----
layer_normalization_12 (Laye (None, 256)          512
-----
dense_13 (Dense)           (None, 10)            2570
=====
Total params: 14,237,450
Trainable params: 14,237,450
Non-trainable params: 0
-----

```

```

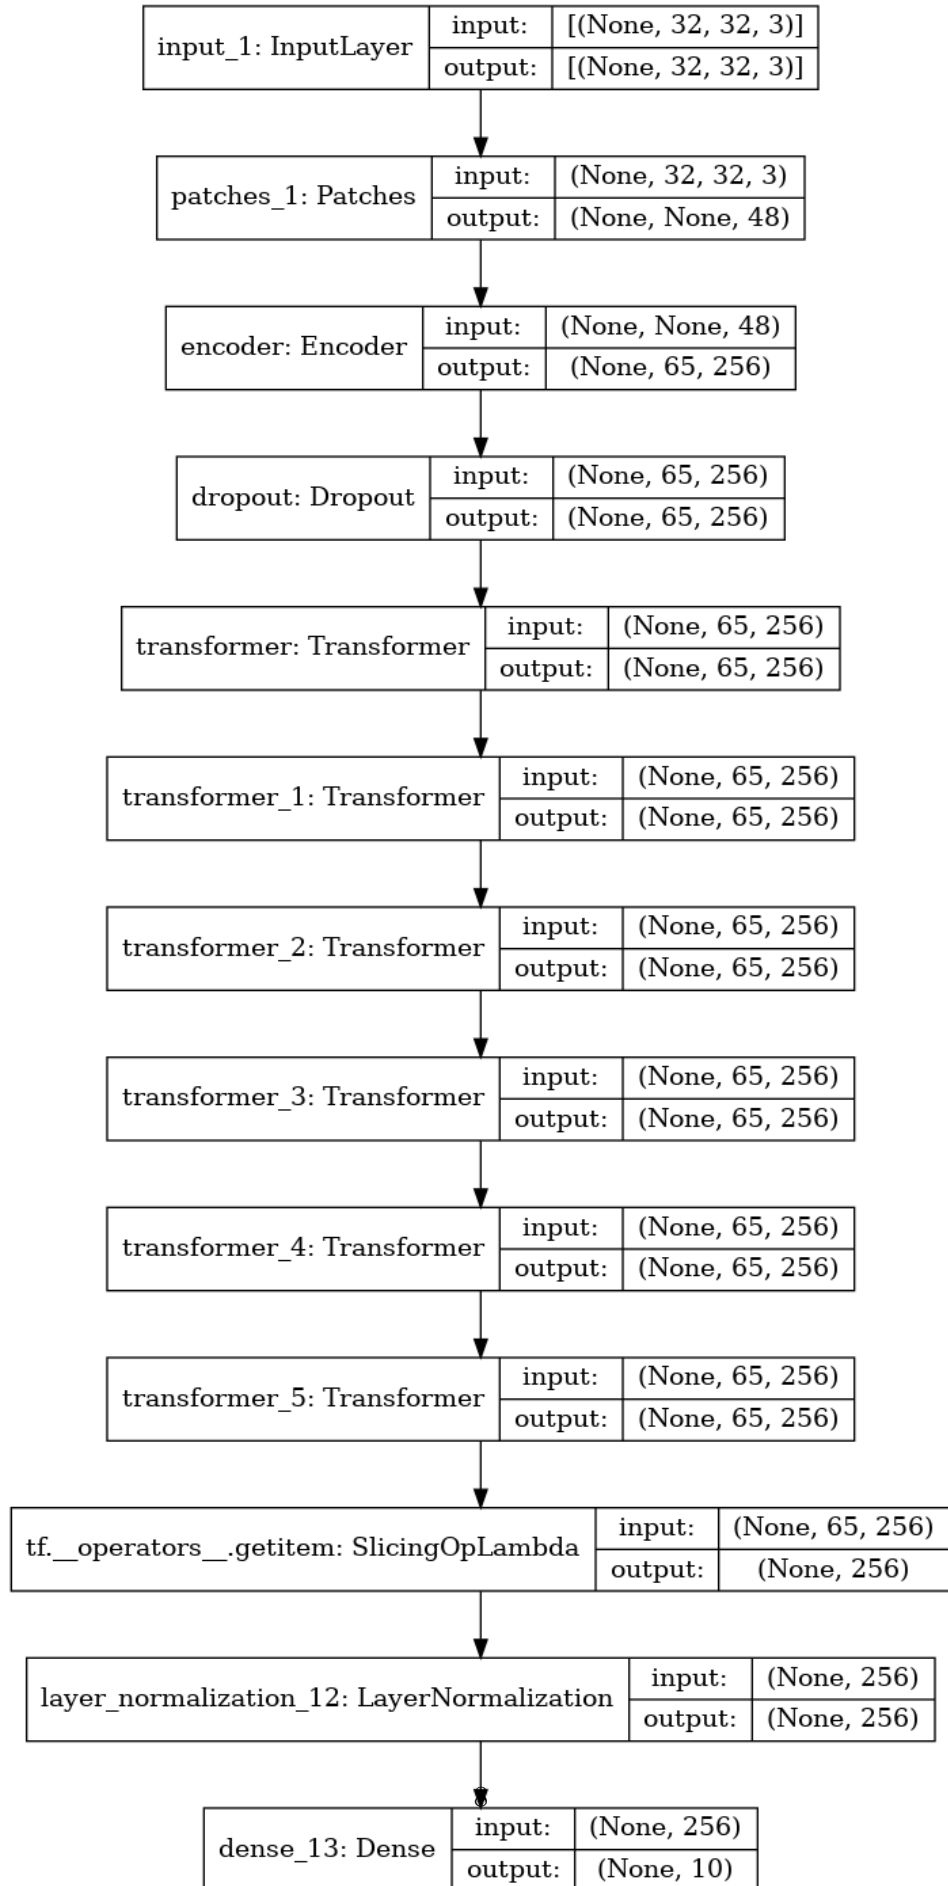
[ ]: tf.keras.utils.plot_model(vit, to_file='./transformer.png', show_shapes=True,
    ↪ expand_nested=True)

```

```

[ ]:

```



Niestety nie umiem rozwinąć wizualizacji bloków transformer i encoder, ale wydaje mi się, że powinny być dobrze.

```
[ ]: def scheduler(epoch, lr):  
    if epoch == 100 or epoch == 150:  
        return lr * tf.constant(0.1)  
    else:  
        return lr  
  
    cb = tf.keras.callbacks.LearningRateScheduler(scheduler)  
  
[ ]: optimizer = tf.keras.optimizers.Adam(learning_rate=0.00001)  
vit.compile(  
    optimizer=optimizer,  
    loss=tf.keras.losses.CategoricalCrossentropy(),  
    metrics=['accuracy']  
)  
round(vit.optimizer.lr.numpy(), 5)
```

```
[ ]: 1e-05
```

```
[ ]: history = vit.fit(  
    x=train_images,  
    y=train_labels,  
    batch_size=128,  
    epochs=160,  
    callbacks=[cb],  
    validation_data=(test_images, test_labels)  
)  
round(vit.optimizer.lr.numpy(), 5)
```

Epoch 1/160

2022-12-02 14:23:03.208208: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369]
Loaded cuDNN version 8005

391/391 [=====] - 109s 260ms/step - loss: 2.3080 -
accuracy: 0.1619 - val_loss: 2.0736 - val_accuracy: 0.2101

Epoch 2/160

391/391 [=====] - 101s 258ms/step - loss: 2.0804 -
accuracy: 0.2134 - val_loss: 1.9628 - val_accuracy: 0.2706

Epoch 3/160

391/391 [=====] - 101s 258ms/step - loss: 1.9682 -
accuracy: 0.2552 - val_loss: 1.8661 - val_accuracy: 0.2949

Epoch 4/160

391/391 [=====] - 101s 258ms/step - loss: 1.8961 -
accuracy: 0.2779 - val_loss: 1.8080 - val_accuracy: 0.3244

Epoch 5/160
391/391 [=====] - 101s 258ms/step - loss: 1.8415 - accuracy: 0.3012 - val_loss: 1.7726 - val_accuracy: 0.3459

Epoch 6/160
391/391 [=====] - 101s 257ms/step - loss: 1.7885 - accuracy: 0.3231 - val_loss: 1.7151 - val_accuracy: 0.3696

Epoch 7/160
391/391 [=====] - 101s 258ms/step - loss: 1.7409 - accuracy: 0.3450 - val_loss: 1.6820 - val_accuracy: 0.3862

Epoch 8/160
391/391 [=====] - 101s 258ms/step - loss: 1.6965 - accuracy: 0.3665 - val_loss: 1.6370 - val_accuracy: 0.4024

Epoch 9/160
391/391 [=====] - 101s 258ms/step - loss: 1.6539 - accuracy: 0.3876 - val_loss: 1.5997 - val_accuracy: 0.4186

Epoch 10/160
391/391 [=====] - 101s 258ms/step - loss: 1.6101 - accuracy: 0.4106 - val_loss: 1.5515 - val_accuracy: 0.4326

Epoch 11/160
391/391 [=====] - 101s 258ms/step - loss: 1.5679 - accuracy: 0.4293 - val_loss: 1.5184 - val_accuracy: 0.4463

Epoch 12/160
391/391 [=====] - 101s 258ms/step - loss: 1.5330 - accuracy: 0.4423 - val_loss: 1.4945 - val_accuracy: 0.4589

Epoch 13/160
391/391 [=====] - 101s 258ms/step - loss: 1.4989 - accuracy: 0.4576 - val_loss: 1.4604 - val_accuracy: 0.4738

Epoch 14/160
391/391 [=====] - 101s 258ms/step - loss: 1.4632 - accuracy: 0.4727 - val_loss: 1.4343 - val_accuracy: 0.4763

Epoch 15/160
391/391 [=====] - 101s 258ms/step - loss: 1.4377 - accuracy: 0.4827 - val_loss: 1.4514 - val_accuracy: 0.4749

Epoch 16/160
391/391 [=====] - 101s 258ms/step - loss: 1.4132 - accuracy: 0.4932 - val_loss: 1.4096 - val_accuracy: 0.4925

Epoch 17/160
391/391 [=====] - 101s 257ms/step - loss: 1.3875 - accuracy: 0.5013 - val_loss: 1.3966 - val_accuracy: 0.4988

Epoch 18/160
391/391 [=====] - 101s 257ms/step - loss: 1.3703 - accuracy: 0.5084 - val_loss: 1.3685 - val_accuracy: 0.5089

Epoch 19/160
391/391 [=====] - 100s 256ms/step - loss: 1.3471 - accuracy: 0.5186 - val_loss: 1.3636 - val_accuracy: 0.5136

Epoch 20/160
391/391 [=====] - 100s 257ms/step - loss: 1.3327 - accuracy: 0.5254 - val_loss: 1.3563 - val_accuracy: 0.5131

Epoch 21/160
391/391 [=====] - 100s 257ms/step - loss: 1.3194 - accuracy: 0.5279 - val_loss: 1.3415 - val_accuracy: 0.5179
Epoch 22/160
391/391 [=====] - 100s 256ms/step - loss: 1.3052 - accuracy: 0.5345 - val_loss: 1.3319 - val_accuracy: 0.5213
Epoch 23/160
391/391 [=====] - 100s 256ms/step - loss: 1.2883 - accuracy: 0.5385 - val_loss: 1.3307 - val_accuracy: 0.5210
Epoch 24/160
391/391 [=====] - 101s 257ms/step - loss: 1.2754 - accuracy: 0.5436 - val_loss: 1.3264 - val_accuracy: 0.5253
Epoch 25/160
391/391 [=====] - 101s 258ms/step - loss: 1.2613 - accuracy: 0.5511 - val_loss: 1.3046 - val_accuracy: 0.5306
Epoch 26/160
391/391 [=====] - 100s 257ms/step - loss: 1.2551 - accuracy: 0.5509 - val_loss: 1.2874 - val_accuracy: 0.5386
Epoch 27/160
391/391 [=====] - 100s 257ms/step - loss: 1.2427 - accuracy: 0.5555 - val_loss: 1.2958 - val_accuracy: 0.5348
Epoch 28/160
391/391 [=====] - 100s 257ms/step - loss: 1.2283 - accuracy: 0.5621 - val_loss: 1.2838 - val_accuracy: 0.5392
Epoch 29/160
391/391 [=====] - 100s 257ms/step - loss: 1.2181 - accuracy: 0.5645 - val_loss: 1.2856 - val_accuracy: 0.5370
Epoch 30/160
391/391 [=====] - 100s 256ms/step - loss: 1.2099 - accuracy: 0.5694 - val_loss: 1.2893 - val_accuracy: 0.5380
Epoch 31/160
391/391 [=====] - 100s 256ms/step - loss: 1.2011 - accuracy: 0.5735 - val_loss: 1.2776 - val_accuracy: 0.5431
Epoch 32/160
391/391 [=====] - 100s 256ms/step - loss: 1.1906 - accuracy: 0.5745 - val_loss: 1.2643 - val_accuracy: 0.5477
Epoch 33/160
391/391 [=====] - 100s 257ms/step - loss: 1.1816 - accuracy: 0.5784 - val_loss: 1.2536 - val_accuracy: 0.5506
Epoch 34/160
391/391 [=====] - 101s 258ms/step - loss: 1.1715 - accuracy: 0.5828 - val_loss: 1.2453 - val_accuracy: 0.5557
Epoch 35/160
391/391 [=====] - 100s 257ms/step - loss: 1.1632 - accuracy: 0.5852 - val_loss: 1.2368 - val_accuracy: 0.5591
Epoch 36/160
391/391 [=====] - 100s 257ms/step - loss: 1.1561 - accuracy: 0.5889 - val_loss: 1.2267 - val_accuracy: 0.5607

Epoch 37/160
 391/391 [=====] - 100s 257ms/step - loss: 1.1501 - accuracy: 0.5926 - val_loss: 1.2304 - val_accuracy: 0.5595

Epoch 38/160
 391/391 [=====] - 100s 257ms/step - loss: 1.1429 - accuracy: 0.5925 - val_loss: 1.2300 - val_accuracy: 0.5639

Epoch 39/160
 391/391 [=====] - 100s 257ms/step - loss: 1.1343 - accuracy: 0.5959 - val_loss: 1.2207 - val_accuracy: 0.5672

Epoch 40/160
 391/391 [=====] - 100s 256ms/step - loss: 1.1249 - accuracy: 0.5991 - val_loss: 1.2274 - val_accuracy: 0.5630

Epoch 41/160
 391/391 [=====] - 100s 256ms/step - loss: 1.1199 - accuracy: 0.6009 - val_loss: 1.2143 - val_accuracy: 0.5678

Epoch 42/160
 391/391 [=====] - 100s 257ms/step - loss: 1.1122 - accuracy: 0.6032 - val_loss: 1.2117 - val_accuracy: 0.5681

Epoch 43/160
 391/391 [=====] - 100s 257ms/step - loss: 1.1069 - accuracy: 0.6059 - val_loss: 1.2073 - val_accuracy: 0.5708

Epoch 44/160
 391/391 [=====] - 100s 257ms/step - loss: 1.0983 - accuracy: 0.6069 - val_loss: 1.2176 - val_accuracy: 0.5673

Epoch 45/160
 391/391 [=====] - 100s 256ms/step - loss: 1.0917 - accuracy: 0.6119 - val_loss: 1.2094 - val_accuracy: 0.5726

Epoch 46/160
 391/391 [=====] - 100s 257ms/step - loss: 1.0882 - accuracy: 0.6118 - val_loss: 1.1961 - val_accuracy: 0.5774

Epoch 47/160
 391/391 [=====] - 101s 257ms/step - loss: 1.0776 - accuracy: 0.6156 - val_loss: 1.2013 - val_accuracy: 0.5703

Epoch 48/160
 391/391 [=====] - 101s 258ms/step - loss: 1.0715 - accuracy: 0.6187 - val_loss: 1.1955 - val_accuracy: 0.5762

Epoch 49/160
 391/391 [=====] - 101s 257ms/step - loss: 1.0631 - accuracy: 0.6212 - val_loss: 1.1909 - val_accuracy: 0.5791

Epoch 50/160
 391/391 [=====] - 100s 257ms/step - loss: 1.0612 - accuracy: 0.6207 - val_loss: 1.1798 - val_accuracy: 0.5792

Epoch 51/160
 391/391 [=====] - 100s 257ms/step - loss: 1.0558 - accuracy: 0.6246 - val_loss: 1.1812 - val_accuracy: 0.5807

Epoch 52/160
 391/391 [=====] - 100s 257ms/step - loss: 1.0474 - accuracy: 0.6286 - val_loss: 1.1721 - val_accuracy: 0.5841

Epoch 53/160
391/391 [=====] - 100s 257ms/step - loss: 1.0458 - accuracy: 0.6261 - val_loss: 1.1722 - val_accuracy: 0.5831
Epoch 54/160
391/391 [=====] - 100s 257ms/step - loss: 1.0376 - accuracy: 0.6276 - val_loss: 1.1898 - val_accuracy: 0.5819
Epoch 55/160
391/391 [=====] - 100s 256ms/step - loss: 1.0307 - accuracy: 0.6306 - val_loss: 1.1815 - val_accuracy: 0.5818
Epoch 56/160
391/391 [=====] - 100s 256ms/step - loss: 1.0281 - accuracy: 0.6345 - val_loss: 1.1680 - val_accuracy: 0.5865
Epoch 57/160
391/391 [=====] - 100s 256ms/step - loss: 1.0181 - accuracy: 0.6374 - val_loss: 1.1998 - val_accuracy: 0.5791
Epoch 58/160
391/391 [=====] - 100s 257ms/step - loss: 1.0154 - accuracy: 0.6376 - val_loss: 1.1809 - val_accuracy: 0.5837
Epoch 59/160
391/391 [=====] - 100s 257ms/step - loss: 1.0100 - accuracy: 0.6371 - val_loss: 1.1851 - val_accuracy: 0.5815
Epoch 60/160
391/391 [=====] - 100s 256ms/step - loss: 1.0050 - accuracy: 0.6410 - val_loss: 1.1712 - val_accuracy: 0.5859
Epoch 61/160
391/391 [=====] - 100s 257ms/step - loss: 0.9972 - accuracy: 0.6423 - val_loss: 1.1681 - val_accuracy: 0.5874
Epoch 62/160
391/391 [=====] - 100s 257ms/step - loss: 0.9908 - accuracy: 0.6459 - val_loss: 1.1701 - val_accuracy: 0.5885
Epoch 63/160
391/391 [=====] - 100s 257ms/step - loss: 0.9879 - accuracy: 0.6497 - val_loss: 1.1735 - val_accuracy: 0.5852
Epoch 64/160
391/391 [=====] - 100s 257ms/step - loss: 0.9828 - accuracy: 0.6490 - val_loss: 1.2074 - val_accuracy: 0.5796
Epoch 65/160
391/391 [=====] - 100s 257ms/step - loss: 0.9734 - accuracy: 0.6524 - val_loss: 1.1737 - val_accuracy: 0.5878
Epoch 66/160
391/391 [=====] - 100s 257ms/step - loss: 0.9704 - accuracy: 0.6524 - val_loss: 1.1649 - val_accuracy: 0.5910
Epoch 67/160
391/391 [=====] - 100s 257ms/step - loss: 0.9652 - accuracy: 0.6561 - val_loss: 1.1656 - val_accuracy: 0.5913
Epoch 68/160
391/391 [=====] - 100s 257ms/step - loss: 0.9603 - accuracy: 0.6555 - val_loss: 1.1602 - val_accuracy: 0.5925

Epoch 69/160
391/391 [=====] - 100s 257ms/step - loss: 0.9554 - accuracy: 0.6603 - val_loss: 1.1559 - val_accuracy: 0.5939
Epoch 70/160
391/391 [=====] - 100s 257ms/step - loss: 0.9485 - accuracy: 0.6601 - val_loss: 1.1754 - val_accuracy: 0.5895
Epoch 71/160
391/391 [=====] - 100s 256ms/step - loss: 0.9465 - accuracy: 0.6617 - val_loss: 1.1756 - val_accuracy: 0.5921
Epoch 72/160
391/391 [=====] - 100s 257ms/step - loss: 0.9411 - accuracy: 0.6623 - val_loss: 1.1568 - val_accuracy: 0.5971
Epoch 73/160
391/391 [=====] - 100s 257ms/step - loss: 0.9356 - accuracy: 0.6649 - val_loss: 1.1805 - val_accuracy: 0.5910
Epoch 74/160
391/391 [=====] - 100s 257ms/step - loss: 0.9310 - accuracy: 0.6670 - val_loss: 1.1663 - val_accuracy: 0.5969
Epoch 75/160
391/391 [=====] - 100s 256ms/step - loss: 0.9244 - accuracy: 0.6720 - val_loss: 1.1661 - val_accuracy: 0.5952
Epoch 76/160
391/391 [=====] - 100s 257ms/step - loss: 0.9201 - accuracy: 0.6708 - val_loss: 1.1677 - val_accuracy: 0.5976
Epoch 77/160
391/391 [=====] - 100s 256ms/step - loss: 0.9140 - accuracy: 0.6750 - val_loss: 1.1774 - val_accuracy: 0.5933
Epoch 78/160
391/391 [=====] - 100s 256ms/step - loss: 0.9084 - accuracy: 0.6742 - val_loss: 1.1619 - val_accuracy: 0.6008
Epoch 79/160
391/391 [=====] - 100s 257ms/step - loss: 0.9026 - accuracy: 0.6782 - val_loss: 1.1840 - val_accuracy: 0.5912
Epoch 80/160
391/391 [=====] - 100s 257ms/step - loss: 0.8929 - accuracy: 0.6803 - val_loss: 1.1845 - val_accuracy: 0.5941
Epoch 81/160
391/391 [=====] - 100s 257ms/step - loss: 0.8898 - accuracy: 0.6827 - val_loss: 1.1735 - val_accuracy: 0.5979
Epoch 82/160
391/391 [=====] - 100s 256ms/step - loss: 0.8902 - accuracy: 0.6814 - val_loss: 1.1622 - val_accuracy: 0.5988
Epoch 83/160
391/391 [=====] - 100s 256ms/step - loss: 0.8888 - accuracy: 0.6820 - val_loss: 1.1737 - val_accuracy: 0.5998
Epoch 84/160
391/391 [=====] - 100s 256ms/step - loss: 0.8786 - accuracy: 0.6861 - val_loss: 1.1657 - val_accuracy: 0.5963

Epoch 85/160
391/391 [=====] - 100s 257ms/step - loss: 0.8721 - accuracy: 0.6861 - val_loss: 1.1778 - val_accuracy: 0.5995
Epoch 86/160
391/391 [=====] - 100s 257ms/step - loss: 0.8690 - accuracy: 0.6902 - val_loss: 1.1719 - val_accuracy: 0.6025
Epoch 87/160
391/391 [=====] - 100s 257ms/step - loss: 0.8620 - accuracy: 0.6903 - val_loss: 1.1737 - val_accuracy: 0.5991
Epoch 88/160
391/391 [=====] - 100s 256ms/step - loss: 0.8552 - accuracy: 0.6940 - val_loss: 1.1703 - val_accuracy: 0.5987
Epoch 89/160
391/391 [=====] - 100s 256ms/step - loss: 0.8536 - accuracy: 0.6939 - val_loss: 1.1666 - val_accuracy: 0.6026
Epoch 90/160
391/391 [=====] - 100s 257ms/step - loss: 0.8469 - accuracy: 0.6977 - val_loss: 1.1771 - val_accuracy: 0.5966
Epoch 91/160
391/391 [=====] - 100s 256ms/step - loss: 0.8399 - accuracy: 0.6984 - val_loss: 1.1792 - val_accuracy: 0.6020
Epoch 92/160
391/391 [=====] - 100s 256ms/step - loss: 0.8335 - accuracy: 0.7041 - val_loss: 1.1722 - val_accuracy: 0.6012
Epoch 93/160
391/391 [=====] - 100s 256ms/step - loss: 0.8327 - accuracy: 0.7010 - val_loss: 1.1831 - val_accuracy: 0.5999
Epoch 94/160
391/391 [=====] - 100s 256ms/step - loss: 0.8247 - accuracy: 0.7027 - val_loss: 1.1893 - val_accuracy: 0.5979
Epoch 95/160
391/391 [=====] - 100s 257ms/step - loss: 0.8221 - accuracy: 0.7060 - val_loss: 1.1883 - val_accuracy: 0.5965
Epoch 96/160
391/391 [=====] - 100s 256ms/step - loss: 0.8172 - accuracy: 0.7067 - val_loss: 1.1991 - val_accuracy: 0.5977
Epoch 97/160
391/391 [=====] - 100s 256ms/step - loss: 0.8096 - accuracy: 0.7116 - val_loss: 1.1856 - val_accuracy: 0.6005
Epoch 98/160
391/391 [=====] - 100s 256ms/step - loss: 0.8077 - accuracy: 0.7106 - val_loss: 1.1916 - val_accuracy: 0.6011
Epoch 99/160
391/391 [=====] - 100s 257ms/step - loss: 0.8026 - accuracy: 0.7116 - val_loss: 1.1918 - val_accuracy: 0.6032
Epoch 100/160
391/391 [=====] - 104s 265ms/step - loss: 0.7951 - accuracy: 0.7149 - val_loss: 1.1936 - val_accuracy: 0.6001

Epoch 101/160
391/391 [=====] - 100s 257ms/step - loss: 0.7674 -
accuracy: 0.7263 - val_loss: 1.1883 - val_accuracy: 0.6035
Epoch 102/160
391/391 [=====] - 100s 256ms/step - loss: 0.7637 -
accuracy: 0.7280 - val_loss: 1.1903 - val_accuracy: 0.6023
Epoch 103/160
391/391 [=====] - 100s 257ms/step - loss: 0.7609 -
accuracy: 0.7281 - val_loss: 1.1911 - val_accuracy: 0.6063
Epoch 104/160
391/391 [=====] - 101s 257ms/step - loss: 0.7621 -
accuracy: 0.7284 - val_loss: 1.1924 - val_accuracy: 0.6058
Epoch 105/160
391/391 [=====] - 101s 258ms/step - loss: 0.7597 -
accuracy: 0.7284 - val_loss: 1.1985 - val_accuracy: 0.6045
Epoch 106/160
391/391 [=====] - 101s 257ms/step - loss: 0.7564 -
accuracy: 0.7305 - val_loss: 1.1914 - val_accuracy: 0.6060
Epoch 107/160
391/391 [=====] - 101s 257ms/step - loss: 0.7582 -
accuracy: 0.7282 - val_loss: 1.1929 - val_accuracy: 0.6053
Epoch 108/160
391/391 [=====] - 101s 257ms/step - loss: 0.7565 -
accuracy: 0.7292 - val_loss: 1.1963 - val_accuracy: 0.6051
Epoch 109/160
391/391 [=====] - 101s 257ms/step - loss: 0.7550 -
accuracy: 0.7303 - val_loss: 1.1955 - val_accuracy: 0.6047
Epoch 110/160
391/391 [=====] - 101s 258ms/step - loss: 0.7559 -
accuracy: 0.7276 - val_loss: 1.1996 - val_accuracy: 0.6056
Epoch 111/160
391/391 [=====] - 101s 257ms/step - loss: 0.7545 -
accuracy: 0.7291 - val_loss: 1.1966 - val_accuracy: 0.6033
Epoch 112/160
391/391 [=====] - 101s 257ms/step - loss: 0.7550 -
accuracy: 0.7278 - val_loss: 1.1923 - val_accuracy: 0.6040
Epoch 113/160
391/391 [=====] - 101s 257ms/step - loss: 0.7527 -
accuracy: 0.7301 - val_loss: 1.1979 - val_accuracy: 0.6054
Epoch 114/160
391/391 [=====] - 101s 257ms/step - loss: 0.7508 -
accuracy: 0.7314 - val_loss: 1.1995 - val_accuracy: 0.6052
Epoch 115/160
391/391 [=====] - 101s 257ms/step - loss: 0.7498 -
accuracy: 0.7311 - val_loss: 1.2039 - val_accuracy: 0.6045
Epoch 116/160
391/391 [=====] - 101s 257ms/step - loss: 0.7551 -
accuracy: 0.7278 - val_loss: 1.1973 - val_accuracy: 0.6033

Epoch 117/160
391/391 [=====] - 101s 257ms/step - loss: 0.7526 - accuracy: 0.7310 - val_loss: 1.2064 - val_accuracy: 0.6051
Epoch 118/160
391/391 [=====] - 101s 258ms/step - loss: 0.7487 - accuracy: 0.7317 - val_loss: 1.1983 - val_accuracy: 0.6049
Epoch 119/160
391/391 [=====] - 101s 257ms/step - loss: 0.7492 - accuracy: 0.7314 - val_loss: 1.1990 - val_accuracy: 0.6064
Epoch 120/160
391/391 [=====] - 101s 257ms/step - loss: 0.7489 - accuracy: 0.7310 - val_loss: 1.1975 - val_accuracy: 0.6051
Epoch 121/160
391/391 [=====] - 101s 257ms/step - loss: 0.7455 - accuracy: 0.7332 - val_loss: 1.2065 - val_accuracy: 0.6044
Epoch 122/160
391/391 [=====] - 101s 257ms/step - loss: 0.7466 - accuracy: 0.7329 - val_loss: 1.1993 - val_accuracy: 0.6062
Epoch 123/160
391/391 [=====] - 101s 257ms/step - loss: 0.7492 - accuracy: 0.7317 - val_loss: 1.2004 - val_accuracy: 0.6060
Epoch 124/160
391/391 [=====] - 100s 257ms/step - loss: 0.7469 - accuracy: 0.7326 - val_loss: 1.2022 - val_accuracy: 0.6052
Epoch 125/160
391/391 [=====] - 101s 257ms/step - loss: 0.7469 - accuracy: 0.7339 - val_loss: 1.2035 - val_accuracy: 0.6033
Epoch 126/160
391/391 [=====] - 101s 257ms/step - loss: 0.7452 - accuracy: 0.7337 - val_loss: 1.2045 - val_accuracy: 0.6045
Epoch 127/160
391/391 [=====] - 101s 257ms/step - loss: 0.7448 - accuracy: 0.7327 - val_loss: 1.2033 - val_accuracy: 0.6039
Epoch 128/160
391/391 [=====] - 101s 257ms/step - loss: 0.7442 - accuracy: 0.7323 - val_loss: 1.2096 - val_accuracy: 0.6040
Epoch 129/160
391/391 [=====] - 101s 257ms/step - loss: 0.7407 - accuracy: 0.7337 - val_loss: 1.2100 - val_accuracy: 0.6047
Epoch 130/160
391/391 [=====] - 101s 257ms/step - loss: 0.7426 - accuracy: 0.7329 - val_loss: 1.2070 - val_accuracy: 0.6023
Epoch 131/160
391/391 [=====] - 101s 257ms/step - loss: 0.7443 - accuracy: 0.7328 - val_loss: 1.2046 - val_accuracy: 0.6042
Epoch 132/160
391/391 [=====] - 101s 258ms/step - loss: 0.7358 - accuracy: 0.7368 - val_loss: 1.2073 - val_accuracy: 0.6012

Epoch 133/160
391/391 [=====] - 101s 257ms/step - loss: 0.7401 -
accuracy: 0.7334 - val_loss: 1.2089 - val_accuracy: 0.6045
Epoch 134/160
391/391 [=====] - 101s 258ms/step - loss: 0.7418 -
accuracy: 0.7350 - val_loss: 1.2073 - val_accuracy: 0.6033
Epoch 135/160
391/391 [=====] - 101s 258ms/step - loss: 0.7384 -
accuracy: 0.7361 - val_loss: 1.2106 - val_accuracy: 0.6043
Epoch 136/160
391/391 [=====] - 101s 257ms/step - loss: 0.7372 -
accuracy: 0.7347 - val_loss: 1.2083 - val_accuracy: 0.6042
Epoch 137/160
391/391 [=====] - 101s 257ms/step - loss: 0.7396 -
accuracy: 0.7353 - val_loss: 1.2124 - val_accuracy: 0.6042
Epoch 138/160
391/391 [=====] - 101s 257ms/step - loss: 0.7389 -
accuracy: 0.7359 - val_loss: 1.2079 - val_accuracy: 0.6044
Epoch 139/160
391/391 [=====] - 101s 257ms/step - loss: 0.7384 -
accuracy: 0.7355 - val_loss: 1.2044 - val_accuracy: 0.6037
Epoch 140/160
391/391 [=====] - 101s 257ms/step - loss: 0.7370 -
accuracy: 0.7354 - val_loss: 1.2070 - val_accuracy: 0.6056
Epoch 141/160
391/391 [=====] - 101s 257ms/step - loss: 0.7363 -
accuracy: 0.7351 - val_loss: 1.2114 - val_accuracy: 0.6061
Epoch 142/160
391/391 [=====] - 101s 257ms/step - loss: 0.7329 -
accuracy: 0.7395 - val_loss: 1.2129 - val_accuracy: 0.6044
Epoch 143/160
391/391 [=====] - 101s 257ms/step - loss: 0.7344 -
accuracy: 0.7363 - val_loss: 1.2079 - val_accuracy: 0.6032
Epoch 144/160
391/391 [=====] - 101s 257ms/step - loss: 0.7362 -
accuracy: 0.7391 - val_loss: 1.2150 - val_accuracy: 0.6038
Epoch 145/160
391/391 [=====] - 101s 257ms/step - loss: 0.7339 -
accuracy: 0.7377 - val_loss: 1.2094 - val_accuracy: 0.6051
Epoch 146/160
391/391 [=====] - 101s 257ms/step - loss: 0.7331 -
accuracy: 0.7354 - val_loss: 1.2124 - val_accuracy: 0.6043
Epoch 147/160
391/391 [=====] - 100s 257ms/step - loss: 0.7328 -
accuracy: 0.7367 - val_loss: 1.2087 - val_accuracy: 0.6054
Epoch 148/160
391/391 [=====] - 101s 257ms/step - loss: 0.7362 -
accuracy: 0.7354 - val_loss: 1.2148 - val_accuracy: 0.6018

```

Epoch 149/160
391/391 [=====] - 101s 257ms/step - loss: 0.7314 -
accuracy: 0.7370 - val_loss: 1.2113 - val_accuracy: 0.6051
Epoch 150/160
391/391 [=====] - 101s 257ms/step - loss: 0.7326 -
accuracy: 0.7363 - val_loss: 1.2142 - val_accuracy: 0.6047
Epoch 151/160
391/391 [=====] - 101s 257ms/step - loss: 0.7292 -
accuracy: 0.7394 - val_loss: 1.2122 - val_accuracy: 0.6044
Epoch 152/160
391/391 [=====] - 100s 257ms/step - loss: 0.7269 -
accuracy: 0.7409 - val_loss: 1.2112 - val_accuracy: 0.6046
Epoch 153/160
391/391 [=====] - 101s 257ms/step - loss: 0.7274 -
accuracy: 0.7382 - val_loss: 1.2113 - val_accuracy: 0.6053
Epoch 154/160
391/391 [=====] - 101s 257ms/step - loss: 0.7274 -
accuracy: 0.7403 - val_loss: 1.2118 - val_accuracy: 0.6048
Epoch 155/160
391/391 [=====] - 101s 257ms/step - loss: 0.7257 -
accuracy: 0.7390 - val_loss: 1.2122 - val_accuracy: 0.6045
Epoch 156/160
391/391 [=====] - 101s 257ms/step - loss: 0.7245 -
accuracy: 0.7397 - val_loss: 1.2124 - val_accuracy: 0.6055
Epoch 157/160
391/391 [=====] - 101s 257ms/step - loss: 0.7290 -
accuracy: 0.7401 - val_loss: 1.2122 - val_accuracy: 0.6048
Epoch 158/160
391/391 [=====] - 101s 257ms/step - loss: 0.7255 -
accuracy: 0.7384 - val_loss: 1.2124 - val_accuracy: 0.6047
Epoch 159/160
391/391 [=====] - 101s 258ms/step - loss: 0.7288 -
accuracy: 0.7396 - val_loss: 1.2126 - val_accuracy: 0.6052
Epoch 160/160
391/391 [=====] - 100s 257ms/step - loss: 0.7271 -
accuracy: 0.7409 - val_loss: 1.2126 - val_accuracy: 0.6049

```

```
[ ]: 0.0
```

```
[ ]: vit.save_weights("./ViTAdam.h5")
```

```
[ ]: vit.optimizer.lr.numpy()
```

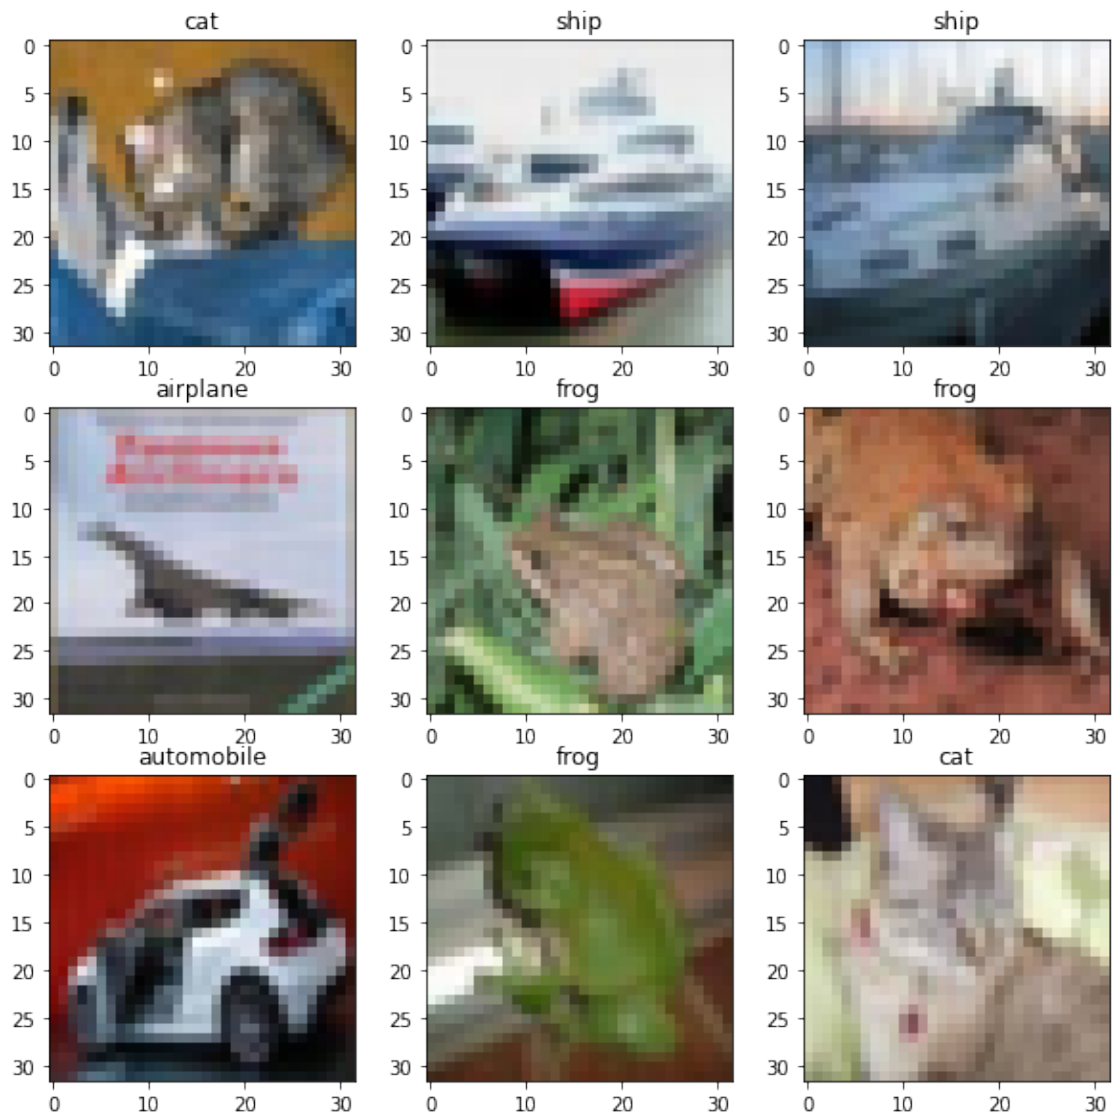
```
[ ]: 1e-07
```

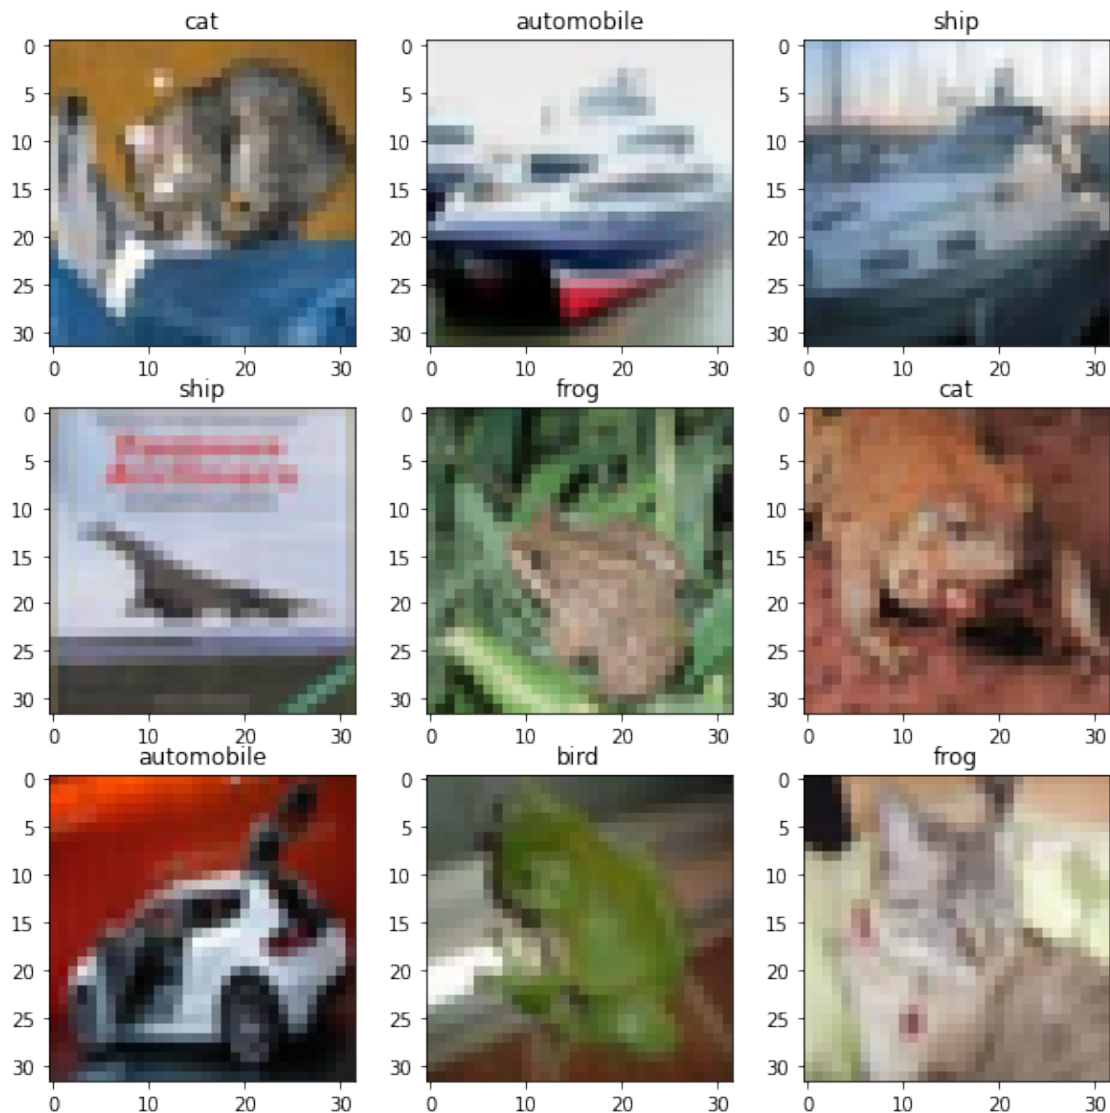
Co ciekawe learning rate na koniec wynosi 1e-07, zaczynaliśmy od 1e-05, więc scheduler zadziałał dobrze (choć myślałem, że adam sam dostosowuje learning rate trochę bardziej?).

Model nauczył się całkiem nieźle, choć niestety dalej nie jakoś wybitnie i na zbiorze testowym

osiąga aż/tylko 60% accuracy.

```
[ ]: def peak_ds(ds, labels):  
    plt.figure(figsize=(10,10))  
    for ix, image in enumerate(test_images[:9]):  
        plt.subplot(3,3, ix+1)  
        plt.imshow(image)  
        plt.title(classes[tf.math.argmax(labels[ix])])  
    plt.show()  
  
peak_ds(test_images[:9], test_labels[:9])  
peak_ds(test_images[:9], vit.predict(test_images[:9]))
```





Jak zobaczymy co zwraca nasz model to można się nawet troszkę zaśmiać :-)) aczkolwiek niektóre wyniki są dobre.

Próbowałem zrobić Attention Rollout, ale niestety przez to w jaki sposób zdefiniowałem model - przy użyciu functional API, i class'y Layer stworzyłem warstwę transformer i nie umiem z niej wyciągnąć informacji o output'cie multihead attention...

Nie mam pomysłu jak to łatwo poprawić, a nie mam też czasu uczyć nową sieć.

Z tego co rozumiem powinienem zaaplikować jakąś funkcję (mean, min) na output'cie multihead attention, tak żeby zebrać razem outputy wszystkich "tokenów" oraz tego trochę nie rozumiem, ale dodajemy macierz identyczności, żeby zasymulować skip connection?. Następnie takie outputy mnożymy między siebie.

Następnie spróbuję wykorzystać AdamW jako optimizer.

(Tak naprawdę najpierw przetestowałem AdamW i wyszedł gorzej, ale nie mam czasu szukać lep-

szych parametrów i zostawiłem AdamW na koniec)

```
[ ]: optimizer = tfa.optimizers.AdamW(
    learning_rate=0.00001,
    weight_decay=0.00004
)
# optimizer = tf.keras.optimizers.Adam(learning_rate=0.00001)
vit.compile(
    optimizer=optimizer,
    loss=tf.keras.losses.CategoricalCrossentropy(),
    metrics=['accuracy']
)
round(vit.optimizer.lr.numpy(), 5)
```

```
[ ]: 1e-05
```

```
[ ]: history = vit.fit(
    x=train_images,
    y=train_labels,
    batch_size=128,
    epochs=160,
    callbacks=[cb],
    validation_data=(test_images, test_labels)
)
round(vit.optimizer.lr.numpy(), 5)
```

Epoch 1/160

2022-12-01 18:51:07.180023: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369]
Loaded cuDNN version 8005

391/391 [=====] - 110s 262ms/step - loss: 2.3118 -
accuracy: 0.1584 - val_loss: 2.0603 - val_accuracy: 0.2379

Epoch 2/160

391/391 [=====] - 101s 259ms/step - loss: 2.0671 -
accuracy: 0.2216 - val_loss: 1.9248 - val_accuracy: 0.2792

Epoch 3/160

391/391 [=====] - 101s 260ms/step - loss: 1.9531 -
accuracy: 0.2556 - val_loss: 1.8527 - val_accuracy: 0.3038

Epoch 4/160

391/391 [=====] - 101s 260ms/step - loss: 1.8719 -
accuracy: 0.2837 - val_loss: 1.8024 - val_accuracy: 0.3224

Epoch 5/160

391/391 [=====] - 101s 259ms/step - loss: 1.8198 -
accuracy: 0.3033 - val_loss: 1.7838 - val_accuracy: 0.3297

Epoch 6/160

391/391 [=====] - 101s 259ms/step - loss: 1.7735 -
accuracy: 0.3261 - val_loss: 1.7145 - val_accuracy: 0.3595

Epoch 7/160

391/391 [=====] - 102s 260ms/step - loss: 1.7342 - accuracy: 0.3451 - val_loss: 1.6734 - val_accuracy: 0.3865
Epoch 8/160
391/391 [=====] - 101s 259ms/step - loss: 1.6926 - accuracy: 0.3650 - val_loss: 1.6385 - val_accuracy: 0.4002
Epoch 9/160
391/391 [=====] - 101s 259ms/step - loss: 1.6508 - accuracy: 0.3865 - val_loss: 1.6215 - val_accuracy: 0.4019
Epoch 10/160
391/391 [=====] - 101s 259ms/step - loss: 1.6114 - accuracy: 0.4062 - val_loss: 1.5779 - val_accuracy: 0.4214
Epoch 11/160
391/391 [=====] - 101s 259ms/step - loss: 1.5627 - accuracy: 0.4255 - val_loss: 1.5198 - val_accuracy: 0.4455
Epoch 12/160
391/391 [=====] - 101s 259ms/step - loss: 1.5210 - accuracy: 0.4446 - val_loss: 1.4730 - val_accuracy: 0.4652
Epoch 13/160
391/391 [=====] - 101s 259ms/step - loss: 1.4819 - accuracy: 0.4612 - val_loss: 1.4510 - val_accuracy: 0.4717
Epoch 14/160
391/391 [=====] - 101s 259ms/step - loss: 1.4493 - accuracy: 0.4733 - val_loss: 1.4239 - val_accuracy: 0.4843
Epoch 15/160
391/391 [=====] - 101s 259ms/step - loss: 1.4197 - accuracy: 0.4882 - val_loss: 1.4184 - val_accuracy: 0.4823
Epoch 16/160
391/391 [=====] - 101s 259ms/step - loss: 1.3976 - accuracy: 0.4943 - val_loss: 1.3853 - val_accuracy: 0.5033
Epoch 17/160
391/391 [=====] - 101s 259ms/step - loss: 1.3775 - accuracy: 0.5018 - val_loss: 1.3794 - val_accuracy: 0.5086
Epoch 18/160
391/391 [=====] - 101s 259ms/step - loss: 1.3615 - accuracy: 0.5120 - val_loss: 1.3577 - val_accuracy: 0.5186
Epoch 19/160
391/391 [=====] - 101s 259ms/step - loss: 1.3454 - accuracy: 0.5187 - val_loss: 1.3498 - val_accuracy: 0.5178
Epoch 20/160
391/391 [=====] - 101s 259ms/step - loss: 1.3335 - accuracy: 0.5230 - val_loss: 1.3500 - val_accuracy: 0.5171
Epoch 21/160
391/391 [=====] - 101s 259ms/step - loss: 1.3222 - accuracy: 0.5257 - val_loss: 1.3520 - val_accuracy: 0.5195
Epoch 22/160
391/391 [=====] - 101s 259ms/step - loss: 1.3119 - accuracy: 0.5316 - val_loss: 1.3214 - val_accuracy: 0.5320
Epoch 23/160

391/391 [=====] - 101s 259ms/step - loss: 1.3007 -
 accuracy: 0.5340 - val_loss: 1.3351 - val_accuracy: 0.5259
 Epoch 24/160
 391/391 [=====] - 105s 267ms/step - loss: 1.2910 -
 accuracy: 0.5356 - val_loss: 1.3181 - val_accuracy: 0.5315
 Epoch 25/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2876 -
 accuracy: 0.5380 - val_loss: 1.3314 - val_accuracy: 0.5285
 Epoch 26/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2778 -
 accuracy: 0.5427 - val_loss: 1.3111 - val_accuracy: 0.5370
 Epoch 27/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2708 -
 accuracy: 0.5458 - val_loss: 1.3203 - val_accuracy: 0.5354
 Epoch 28/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2629 -
 accuracy: 0.5500 - val_loss: 1.2969 - val_accuracy: 0.5450
 Epoch 29/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2615 -
 accuracy: 0.5500 - val_loss: 1.3068 - val_accuracy: 0.5379
 Epoch 30/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2551 -
 accuracy: 0.5535 - val_loss: 1.2970 - val_accuracy: 0.5413
 Epoch 31/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2492 -
 accuracy: 0.5538 - val_loss: 1.3094 - val_accuracy: 0.5394
 Epoch 32/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2444 -
 accuracy: 0.5561 - val_loss: 1.3048 - val_accuracy: 0.5366
 Epoch 33/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2370 -
 accuracy: 0.5605 - val_loss: 1.2850 - val_accuracy: 0.5495
 Epoch 34/160
 391/391 [=====] - 101s 260ms/step - loss: 1.2351 -
 accuracy: 0.5610 - val_loss: 1.3065 - val_accuracy: 0.5410
 Epoch 35/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2325 -
 accuracy: 0.5639 - val_loss: 1.3028 - val_accuracy: 0.5433
 Epoch 36/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2272 -
 accuracy: 0.5649 - val_loss: 1.2945 - val_accuracy: 0.5451
 Epoch 37/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2272 -
 accuracy: 0.5648 - val_loss: 1.2841 - val_accuracy: 0.5462
 Epoch 38/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2235 -
 accuracy: 0.5663 - val_loss: 1.2758 - val_accuracy: 0.5517
 Epoch 39/160

391/391 [=====] - 101s 259ms/step - loss: 1.2230 -
accuracy: 0.5664 - val_loss: 1.2807 - val_accuracy: 0.5498
Epoch 40/160
391/391 [=====] - 101s 259ms/step - loss: 1.2173 -
accuracy: 0.5683 - val_loss: 1.3007 - val_accuracy: 0.5432
Epoch 41/160
391/391 [=====] - 101s 259ms/step - loss: 1.2144 -
accuracy: 0.5720 - val_loss: 1.2779 - val_accuracy: 0.5525
Epoch 42/160
391/391 [=====] - 101s 259ms/step - loss: 1.2130 -
accuracy: 0.5725 - val_loss: 1.2786 - val_accuracy: 0.5509
Epoch 43/160
391/391 [=====] - 101s 259ms/step - loss: 1.2112 -
accuracy: 0.5744 - val_loss: 1.2711 - val_accuracy: 0.5566
Epoch 44/160
391/391 [=====] - 101s 259ms/step - loss: 1.2116 -
accuracy: 0.5733 - val_loss: 1.2683 - val_accuracy: 0.5534
Epoch 45/160
391/391 [=====] - 101s 259ms/step - loss: 1.2108 -
accuracy: 0.5744 - val_loss: 1.2754 - val_accuracy: 0.5501
Epoch 46/160
391/391 [=====] - 101s 259ms/step - loss: 1.2089 -
accuracy: 0.5749 - val_loss: 1.2905 - val_accuracy: 0.5488
Epoch 47/160
391/391 [=====] - 101s 259ms/step - loss: 1.2038 -
accuracy: 0.5799 - val_loss: 1.2913 - val_accuracy: 0.5448
Epoch 48/160
391/391 [=====] - 102s 260ms/step - loss: 1.2040 -
accuracy: 0.5780 - val_loss: 1.2606 - val_accuracy: 0.5542
Epoch 49/160
391/391 [=====] - 101s 259ms/step - loss: 1.2042 -
accuracy: 0.5791 - val_loss: 1.2589 - val_accuracy: 0.5574
Epoch 50/160
391/391 [=====] - 101s 259ms/step - loss: 1.2044 -
accuracy: 0.5796 - val_loss: 1.2813 - val_accuracy: 0.5473
Epoch 51/160
391/391 [=====] - 101s 259ms/step - loss: 1.2017 -
accuracy: 0.5803 - val_loss: 1.2654 - val_accuracy: 0.5544
Epoch 52/160
391/391 [=====] - 101s 259ms/step - loss: 1.2003 -
accuracy: 0.5791 - val_loss: 1.2607 - val_accuracy: 0.5568
Epoch 53/160
391/391 [=====] - 101s 259ms/step - loss: 1.1992 -
accuracy: 0.5801 - val_loss: 1.2784 - val_accuracy: 0.5526
Epoch 54/160
391/391 [=====] - 101s 259ms/step - loss: 1.2006 -
accuracy: 0.5789 - val_loss: 1.2518 - val_accuracy: 0.5623
Epoch 55/160

391/391 [=====] - 101s 259ms/step - loss: 1.2012 -
accuracy: 0.5801 - val_loss: 1.2543 - val_accuracy: 0.5587
Epoch 56/160
391/391 [=====] - 101s 259ms/step - loss: 1.1997 -
accuracy: 0.5830 - val_loss: 1.2691 - val_accuracy: 0.5546
Epoch 57/160
391/391 [=====] - 101s 259ms/step - loss: 1.1990 -
accuracy: 0.5816 - val_loss: 1.2649 - val_accuracy: 0.5523
Epoch 58/160
391/391 [=====] - 101s 259ms/step - loss: 1.1983 -
accuracy: 0.5856 - val_loss: 1.2554 - val_accuracy: 0.5580
Epoch 59/160
391/391 [=====] - 101s 259ms/step - loss: 1.1990 -
accuracy: 0.5847 - val_loss: 1.2544 - val_accuracy: 0.5588
Epoch 60/160
391/391 [=====] - 101s 259ms/step - loss: 1.1958 -
accuracy: 0.5845 - val_loss: 1.2607 - val_accuracy: 0.5552
Epoch 61/160
391/391 [=====] - 101s 259ms/step - loss: 1.1958 -
accuracy: 0.5856 - val_loss: 1.2650 - val_accuracy: 0.5535
Epoch 62/160
391/391 [=====] - 101s 259ms/step - loss: 1.1978 -
accuracy: 0.5839 - val_loss: 1.2557 - val_accuracy: 0.5521
Epoch 63/160
391/391 [=====] - 101s 259ms/step - loss: 1.1939 -
accuracy: 0.5879 - val_loss: 1.2759 - val_accuracy: 0.5538
Epoch 64/160
391/391 [=====] - 101s 259ms/step - loss: 1.1958 -
accuracy: 0.5865 - val_loss: 1.2460 - val_accuracy: 0.5645
Epoch 65/160
391/391 [=====] - 101s 259ms/step - loss: 1.1953 -
accuracy: 0.5871 - val_loss: 1.2598 - val_accuracy: 0.5534
Epoch 66/160
391/391 [=====] - 101s 259ms/step - loss: 1.1966 -
accuracy: 0.5844 - val_loss: 1.2785 - val_accuracy: 0.5519
Epoch 67/160
391/391 [=====] - 101s 259ms/step - loss: 1.1970 -
accuracy: 0.5854 - val_loss: 1.2596 - val_accuracy: 0.5565
Epoch 68/160
391/391 [=====] - 101s 259ms/step - loss: 1.1929 -
accuracy: 0.5873 - val_loss: 1.2440 - val_accuracy: 0.5640
Epoch 69/160
391/391 [=====] - 101s 259ms/step - loss: 1.1948 -
accuracy: 0.5859 - val_loss: 1.2501 - val_accuracy: 0.5621
Epoch 70/160
391/391 [=====] - 101s 259ms/step - loss: 1.1942 -
accuracy: 0.5870 - val_loss: 1.2593 - val_accuracy: 0.5517
Epoch 71/160

391/391 [=====] - 101s 259ms/step - loss: 1.1945 -
accuracy: 0.5870 - val_loss: 1.2531 - val_accuracy: 0.5547
Epoch 72/160

391/391 [=====] - 101s 259ms/step - loss: 1.1935 -
accuracy: 0.5879 - val_loss: 1.2635 - val_accuracy: 0.5527
Epoch 73/160

391/391 [=====] - 101s 259ms/step - loss: 1.1958 -
accuracy: 0.5872 - val_loss: 1.2576 - val_accuracy: 0.5580
Epoch 74/160

391/391 [=====] - 101s 259ms/step - loss: 1.1940 -
accuracy: 0.5880 - val_loss: 1.2639 - val_accuracy: 0.5542
Epoch 75/160

391/391 [=====] - 101s 259ms/step - loss: 1.1930 -
accuracy: 0.5907 - val_loss: 1.2553 - val_accuracy: 0.5596
Epoch 76/160

391/391 [=====] - 101s 259ms/step - loss: 1.1941 -
accuracy: 0.5880 - val_loss: 1.2839 - val_accuracy: 0.5474
Epoch 77/160

391/391 [=====] - 101s 259ms/step - loss: 1.1953 -
accuracy: 0.5872 - val_loss: 1.2487 - val_accuracy: 0.5588
Epoch 78/160

391/391 [=====] - 101s 259ms/step - loss: 1.1940 -
accuracy: 0.5891 - val_loss: 1.3043 - val_accuracy: 0.5341
Epoch 79/160

391/391 [=====] - 101s 259ms/step - loss: 1.1938 -
accuracy: 0.5884 - val_loss: 1.2687 - val_accuracy: 0.5502
Epoch 80/160

391/391 [=====] - 101s 259ms/step - loss: 1.1947 -
accuracy: 0.5909 - val_loss: 1.2759 - val_accuracy: 0.5434
Epoch 81/160

391/391 [=====] - 101s 259ms/step - loss: 1.1921 -
accuracy: 0.5899 - val_loss: 1.3024 - val_accuracy: 0.5359
Epoch 82/160

391/391 [=====] - 101s 259ms/step - loss: 1.1931 -
accuracy: 0.5907 - val_loss: 1.2653 - val_accuracy: 0.5503
Epoch 83/160

391/391 [=====] - 101s 259ms/step - loss: 1.1947 -
accuracy: 0.5887 - val_loss: 1.2612 - val_accuracy: 0.5475
Epoch 84/160

391/391 [=====] - 101s 259ms/step - loss: 1.1929 -
accuracy: 0.5913 - val_loss: 1.2668 - val_accuracy: 0.5530
Epoch 85/160

391/391 [=====] - 101s 259ms/step - loss: 1.1926 -
accuracy: 0.5921 - val_loss: 1.2618 - val_accuracy: 0.5560
Epoch 86/160

391/391 [=====] - 101s 259ms/step - loss: 1.1944 -
accuracy: 0.5888 - val_loss: 1.2432 - val_accuracy: 0.5655
Epoch 87/160

391/391 [=====] - 101s 259ms/step - loss: 1.1922 - accuracy: 0.5900 - val_loss: 1.2538 - val_accuracy: 0.5576
Epoch 88/160
391/391 [=====] - 101s 259ms/step - loss: 1.1935 - accuracy: 0.5906 - val_loss: 1.2427 - val_accuracy: 0.5657
Epoch 89/160
391/391 [=====] - 101s 259ms/step - loss: 1.1923 - accuracy: 0.5899 - val_loss: 1.2390 - val_accuracy: 0.5664
Epoch 90/160
391/391 [=====] - 101s 259ms/step - loss: 1.1886 - accuracy: 0.5895 - val_loss: 1.2578 - val_accuracy: 0.5535
Epoch 91/160
391/391 [=====] - 101s 259ms/step - loss: 1.1937 - accuracy: 0.5902 - val_loss: 1.2474 - val_accuracy: 0.5592
Epoch 92/160
391/391 [=====] - 101s 259ms/step - loss: 1.1913 - accuracy: 0.5920 - val_loss: 1.2507 - val_accuracy: 0.5572
Epoch 93/160
391/391 [=====] - 101s 259ms/step - loss: 1.1921 - accuracy: 0.5902 - val_loss: 1.2568 - val_accuracy: 0.5542
Epoch 94/160
391/391 [=====] - 101s 259ms/step - loss: 1.1909 - accuracy: 0.5922 - val_loss: 1.2408 - val_accuracy: 0.5597
Epoch 95/160
391/391 [=====] - 101s 259ms/step - loss: 1.1935 - accuracy: 0.5923 - val_loss: 1.2505 - val_accuracy: 0.5590
Epoch 96/160
391/391 [=====] - 101s 259ms/step - loss: 1.1895 - accuracy: 0.5903 - val_loss: 1.2377 - val_accuracy: 0.5663
Epoch 97/160
391/391 [=====] - 101s 259ms/step - loss: 1.1899 - accuracy: 0.5920 - val_loss: 1.2530 - val_accuracy: 0.5563
Epoch 98/160
391/391 [=====] - 101s 259ms/step - loss: 1.1929 - accuracy: 0.5908 - val_loss: 1.2647 - val_accuracy: 0.5547
Epoch 99/160
391/391 [=====] - 101s 259ms/step - loss: 1.1929 - accuracy: 0.5900 - val_loss: 1.2413 - val_accuracy: 0.5625
Epoch 100/160
391/391 [=====] - 101s 259ms/step - loss: 1.1953 - accuracy: 0.5899 - val_loss: 1.2585 - val_accuracy: 0.5500
Epoch 101/160
391/391 [=====] - 101s 259ms/step - loss: 1.1792 - accuracy: 0.6028 - val_loss: 1.2432 - val_accuracy: 0.5651
Epoch 102/160
391/391 [=====] - 101s 259ms/step - loss: 1.2032 - accuracy: 0.5988 - val_loss: 1.2631 - val_accuracy: 0.5563
Epoch 103/160

391/391 [=====] - 101s 259ms/step - loss: 1.2340 -
 accuracy: 0.5942 - val_loss: 1.2777 - val_accuracy: 0.5547
 Epoch 104/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2648 -
 accuracy: 0.5888 - val_loss: 1.2978 - val_accuracy: 0.5486
 Epoch 105/160
 391/391 [=====] - 101s 259ms/step - loss: 1.2961 -
 accuracy: 0.5838 - val_loss: 1.3207 - val_accuracy: 0.5445
 Epoch 106/160
 391/391 [=====] - 101s 259ms/step - loss: 1.3270 -
 accuracy: 0.5783 - val_loss: 1.3454 - val_accuracy: 0.5368
 Epoch 107/160
 391/391 [=====] - 101s 259ms/step - loss: 1.3592 -
 accuracy: 0.5696 - val_loss: 1.3610 - val_accuracy: 0.5344
 Epoch 108/160
 391/391 [=====] - 101s 259ms/step - loss: 1.3896 -
 accuracy: 0.5630 - val_loss: 1.3951 - val_accuracy: 0.5236
 Epoch 109/160
 391/391 [=====] - 101s 259ms/step - loss: 1.4195 -
 accuracy: 0.5555 - val_loss: 1.4126 - val_accuracy: 0.5210
 Epoch 110/160
 391/391 [=====] - 101s 259ms/step - loss: 1.4453 -
 accuracy: 0.5468 - val_loss: 1.4330 - val_accuracy: 0.5178
 Epoch 111/160
 391/391 [=====] - 101s 259ms/step - loss: 1.4718 -
 accuracy: 0.5402 - val_loss: 1.4542 - val_accuracy: 0.5112
 Epoch 112/160
 391/391 [=====] - 101s 259ms/step - loss: 1.4968 -
 accuracy: 0.5355 - val_loss: 1.4777 - val_accuracy: 0.5056
 Epoch 113/160
 391/391 [=====] - 101s 259ms/step - loss: 1.5203 -
 accuracy: 0.5249 - val_loss: 1.5005 - val_accuracy: 0.4990
 Epoch 114/160
 391/391 [=====] - 101s 259ms/step - loss: 1.5416 -
 accuracy: 0.5205 - val_loss: 1.5165 - val_accuracy: 0.4943
 Epoch 115/160
 391/391 [=====] - 101s 259ms/step - loss: 1.5629 -
 accuracy: 0.5127 - val_loss: 1.5377 - val_accuracy: 0.4901
 Epoch 116/160
 391/391 [=====] - 101s 259ms/step - loss: 1.5823 -
 accuracy: 0.5049 - val_loss: 1.5546 - val_accuracy: 0.4813
 Epoch 117/160
 391/391 [=====] - 101s 259ms/step - loss: 1.6008 -
 accuracy: 0.5004 - val_loss: 1.5768 - val_accuracy: 0.4748
 Epoch 118/160
 391/391 [=====] - 101s 259ms/step - loss: 1.6193 -
 accuracy: 0.4928 - val_loss: 1.5985 - val_accuracy: 0.4702
 Epoch 119/160

391/391 [=====] - 101s 259ms/step - loss: 1.6354 -
 accuracy: 0.4873 - val_loss: 1.6092 - val_accuracy: 0.4675
 Epoch 120/160
 391/391 [=====] - 101s 259ms/step - loss: 1.6514 -
 accuracy: 0.4817 - val_loss: 1.6254 - val_accuracy: 0.4617
 Epoch 121/160
 391/391 [=====] - 101s 259ms/step - loss: 1.6667 -
 accuracy: 0.4764 - val_loss: 1.6410 - val_accuracy: 0.4570
 Epoch 122/160
 391/391 [=====] - 101s 259ms/step - loss: 1.6806 -
 accuracy: 0.4739 - val_loss: 1.6541 - val_accuracy: 0.4509
 Epoch 123/160
 391/391 [=====] - 101s 259ms/step - loss: 1.6942 -
 accuracy: 0.4664 - val_loss: 1.6733 - val_accuracy: 0.4467
 Epoch 124/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7070 -
 accuracy: 0.4620 - val_loss: 1.6765 - val_accuracy: 0.4462
 Epoch 125/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7194 -
 accuracy: 0.4562 - val_loss: 1.7004 - val_accuracy: 0.4323
 Epoch 126/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7309 -
 accuracy: 0.4524 - val_loss: 1.7071 - val_accuracy: 0.4332
 Epoch 127/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7427 -
 accuracy: 0.4489 - val_loss: 1.7145 - val_accuracy: 0.4337
 Epoch 128/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7535 -
 accuracy: 0.4436 - val_loss: 1.7340 - val_accuracy: 0.4226
 Epoch 129/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7640 -
 accuracy: 0.4377 - val_loss: 1.7455 - val_accuracy: 0.4138
 Epoch 130/160
 391/391 [=====] - 105s 267ms/step - loss: 1.7740 -
 accuracy: 0.4343 - val_loss: 1.7567 - val_accuracy: 0.4074
 Epoch 131/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7830 -
 accuracy: 0.4313 - val_loss: 1.7622 - val_accuracy: 0.4101
 Epoch 132/160
 391/391 [=====] - 101s 259ms/step - loss: 1.7935 -
 accuracy: 0.4264 - val_loss: 1.7777 - val_accuracy: 0.3998
 Epoch 133/160
 391/391 [=====] - 101s 259ms/step - loss: 1.8029 -
 accuracy: 0.4223 - val_loss: 1.7781 - val_accuracy: 0.4082
 Epoch 134/160
 391/391 [=====] - 101s 259ms/step - loss: 1.8113 -
 accuracy: 0.4177 - val_loss: 1.7911 - val_accuracy: 0.4062
 Epoch 135/160

```

391/391 [=====] - 101s 259ms/step - loss: 1.8200 -
accuracy: 0.4164 - val_loss: 1.7947 - val_accuracy: 0.4033
Epoch 136/160
391/391 [=====] - 101s 259ms/step - loss: 1.8286 -
accuracy: 0.4098 - val_loss: 1.8071 - val_accuracy: 0.3984
Epoch 137/160
391/391 [=====] - 101s 259ms/step - loss: 1.8365 -
accuracy: 0.4101 - val_loss: 1.8116 - val_accuracy: 0.3999
Epoch 138/160
391/391 [=====] - 101s 259ms/step - loss: 1.8447 -
accuracy: 0.4050 - val_loss: 1.8235 - val_accuracy: 0.3904
Epoch 139/160
391/391 [=====] - 101s 259ms/step - loss: 1.8524 -
accuracy: 0.4023 - val_loss: 1.8273 - val_accuracy: 0.3898
Epoch 140/160
391/391 [=====] - 101s 259ms/step - loss: 1.8600 -
accuracy: 0.3998 - val_loss: 1.8364 - val_accuracy: 0.3910
Epoch 141/160
391/391 [=====] - 101s 259ms/step - loss: 1.8677 -
accuracy: 0.3947 - val_loss: 1.8435 - val_accuracy: 0.3865
Epoch 142/160
391/391 [=====] - 101s 259ms/step - loss: 1.8751 -
accuracy: 0.3937 - val_loss: 1.8555 - val_accuracy: 0.3734
Epoch 143/160
391/391 [=====] - 101s 259ms/step - loss: 1.8823 -
accuracy: 0.3890 - val_loss: 1.8580 - val_accuracy: 0.3822
Epoch 144/160
391/391 [=====] - 101s 259ms/step - loss: 1.8894 -
accuracy: 0.3874 - val_loss: 1.8652 - val_accuracy: 0.3803
Epoch 145/160
391/391 [=====] - 101s 259ms/step - loss: 1.8963 -
accuracy: 0.3857 - val_loss: 1.8751 - val_accuracy: 0.3728
Epoch 146/160
391/391 [=====] - 101s 259ms/step - loss: 1.9033 -
accuracy: 0.3810 - val_loss: 1.8828 - val_accuracy: 0.3694
Epoch 147/160
391/391 [=====] - 101s 259ms/step - loss: 1.9093 -
accuracy: 0.3787 - val_loss: 1.8919 - val_accuracy: 0.3615
Epoch 148/160
341/391 [=====>...] - ETA: 12s - loss: 1.9164 - accuracy:
0.3742

```

Niestety obliczenia zawiesiły się przy 148 epokach i zawiesiły notebook, ale widać że być może parametry nie były dobrane najlepiej, ponieważ najlepsze accuracy jest w okolicach 0.55 przy 50 epoche, a później zaczęło spadać. Niestety nie mam czasu próbować szukać innych parametrów.

[]: